

**Q1: Explain the primary differences between TensorFlow and PyTorch. When would you choose one over the other?**

Aspect	TensorFlow	PyTorch
<b>Computation Graph</b>	Uses static computation graphs (though TF 2.x allows eager execution).	Uses dynamic computation graphs (define-by-run).
<b>Ease of Debugging</b>	Harder to debug due to static graphs (before TF 2.x).	Easier to debug because graphs are built on the fly, similar to normal Python code.
<b>Deployment</b>	Better for production/deployment (TensorFlow Serving, TensorFlow Lite, TensorFlow.js).	Deployment is possible (TorchServe, ONNX), but less mature.
<b>Community &amp; Ecosystem</b>	Larger ecosystem for production, mobile, and cross-platform.	Strong research community; often the first to receive cutting-edge models.

**When to choose:**

- **PyTorch:** Research, prototyping, experimental models, or when you need flexibility in graph execution.
- **TensorFlow:** Production-ready applications, deployment across devices, and integration into large-scale pipelines.

**Q2: Describe two use cases for Jupyter Notebooks in AI development.**

**1. Exploratory Data Analysis (EDA):**

- Visualizing data distributions, correlations, missing values.
- Iterative testing of preprocessing steps with immediate feedback.

**2. Model Prototyping and Documentation:**

- Quickly build, train, and evaluate ML/DL models.
- Combine code, visualizations, and markdown explanations for reproducibility and sharing with teams.

**Q3: How does spaCy enhance NLP tasks compared to basic Python string operations?**

- **Pre-built NLP Pipelines:** Tokenization, lemmatization, POS tagging, named entity recognition (NER), and dependency parsing.
- **Speed & Accuracy:** Optimized Cython backend for large-scale text, faster than naive Python string operations.

- **Consistency & Reliability:** Handles edge cases (punctuation, contractions, multilingual support) that simple string operations often fail to address.
- **Integration:** Easy interoperability with ML pipelines for downstream tasks like sentiment analysis or text classification.

## 2. Comparative Analysis: Scikit-learn vs. TensorFlow

Feature	Scikit-learn	TensorFlow
<b>Target Applications</b>	Classical ML: regression, classification, clustering, feature engineering.	Deep learning: CNNs, RNNs, Transformers, large-scale neural networks.
<b>Ease of Use for Beginners</b>	Very beginner-friendly: simple API, minimal setup, concise syntax.	Steeper learning curve, especially for beginners in computational graphs and tensors.
<b>Community Support</b>	Mature, large community; many tutorials for classical ML.	Very active research and industry community; extensive documentation and examples for deep learning.

### Summary:

- Use **Scikit-learn** for small-to-medium datasets, classical ML, and rapid prototyping.
- Use **TensorFlow** for complex neural networks, large datasets, and production-ready deployment pipelines.

## PART 3

# 1. Ethical Considerations

When working with **MNIST** (handwritten digits) or **Amazon Reviews** (text data), potential biases can arise:

## Potential Biases

### MNIST:

- **Digit representation bias:** MNIST mostly has digits written by American students in the 1990s. Digits from other regions may be written differently, so the model may misclassify.
- **Class imbalance:** If some digits appear more frequently than others, the model may favor the more common digits.

### Amazon Reviews:

- **Sentiment bias:** Reviews may have cultural or product-category biases. For example, some products may receive more positive reviews because of brand popularity rather than quality.

- **Demographic bias:** Language used in reviews may reflect age, gender, or cultural background, influencing sentiment detection unfairly.
- **Neglecting sarcasm or context:** Rule-based systems may misclassify sarcastic reviews as positive.

## Mitigating Biases

### Tools like TensorFlow Fairness Indicators:

- Can **measure fairness** across groups (e.g., gender, location, or review category).
- Helps **identify if the model treats some groups unfairly**.
- Example: Check if positive/negative sentiment predictions are balanced across product categories.

### spaCy's rule-based systems:

- Can **preprocess text to normalize language**, detect and adjust for bias patterns.
- For example:
  - Remove words that reflect demographic bias.
  - Add rules to detect sarcasm or negation in sentiment analysis.

### Summary:

The key is **awareness and monitoring**—use fairness tools to identify bias and NLP rules to reduce it before model training.

## 2. Troubleshooting Challenge

Imagine a **buggy TensorFlow script** like this:

```
import tensorflow as tf
from tensorflow.keras import layers, models

model = models.Sequential([
    layers.Dense(64, activation='relu', input_shape=(28,28)),
    layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='binary_crossentropy', # ✗ wrong loss
              metrics=['accuracy'])

# x_train shape: (60000,28,28)
# y_train shape: (60000,) # labels 0-9
```

## Problems:

1. **Input shape mismatch:** Dense layer expects 1D input, but MNIST images are (28, 28).
2. **Wrong loss function:** `binary_crossentropy` is for 2-class problems, not 10 classes.
3. **y\_train shape:** Needs one-hot encoding for `categorical_crossentropy`.

## Fixed Code:

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.utils import to_categorical

# Flatten input and one-hot encode labels
x_train = x_train.reshape(-1, 28*28).astype('float32') / 255
y_train = to_categorical(y_train, 10)

x_test = x_test.reshape(-1, 28*28).astype('float32') / 255
y_test = to_categorical(y_test, 10)

model = models.Sequential([
    layers.Dense(64, activation='relu', input_shape=(28*28,)),
    layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5, batch_size=32, validation_data=(x_test,
y_test))
```

## Changes made:

- Flattened input (28\*28, ).
- Changed loss to `categorical_crossentropy`.
- One-hot encoded labels.
- Normalized pixel values.